


1 General Definitions

NIST Cloud: On-demand self-service, Broad network access, Resource pooling (multi-tenant), Rapid elasticity, Measured service (pay-as-you-go). **Cloud Operations:** manage/optimize cloud services & infra. **GitOps:** Git = single source of truth; declarative, versioned CI/CD. **DevOps:** dev+ops; automation, CI/CD, monitoring, agile. **CI:** every push auto-built & tested. **Platform Engineering:** internal self-service dev platforms (IDP); cultural mindset. **DevOps 6 pillars:** Cultural mindset, Automation-first, Continuous Processes, Monitoring/Feedback, Rapid+Reliable Releases, Security.

1.1 DevOps Cycle



Arcs → **stages:**
CI=Code+Build+Test, **CD**=Release+Deploy (+Test), **CDe-livery**=Test+Release (manual prod), **CFeedback**=Operate+Monitor+Plan.
Stages (activities | tools): **Plan** (requirements, roadmap, arch sketch, risk/rollback | Jira, Asana, GitHub Issues, Trello), **Code** (write/refactor, local build/test, peer review | GitHub, GitHub, Bitbucket), **Build** (compile, container images, publish artifacts, versioned packages | GitLab CI, GH Actions, Jenkins, Maven/Gradle, Docker/Podman, Artifactory, Quay), **Test** (unit/integration/E2E, perf/load, sec scans, coverage, quality gates | JUnit/pytest, Selenium/Cypress/Playwright, JMeter/k6, Snyk/Trivy, SonarQube), **Release** (promote artifact, immutable version, release notes, signing | Cosign, GnuPG, GitLab/GH Packages), **Deploy** (pull artifact, apply manifest, strategy, post-deploy validation, rollback | ArgoCD, Flux, Helm, Kustomize, Terraform/OpenTofu, Ansible, Flagger), **Operate** (troubleshoot, scale, patch, secret rotation, drift detect, backups, compliance | PagerDuty, Vault, Velero, HPA/VPA, Falco), **Monitor** (SLI/SLO, dashboards, alerts, synthetic checks, correlate logs/metrics/traces | Prometheus, Grafana, Loki, Jaeger, ELK, Alertmanager).

Delivery vs Deployment: Deployment auto-pushes staging→prod;

Delivery manual gate. **CD Deployment Strategies:**

- **Rolling:** gradual update, minimal downtime
- **Blue-Green:** two envs, fast switch
- **Canary:** small user group first
- **Feature Flag:** deploy, toggle later
- **Dark Launching:** real traffic, hidden from users

1.2 Platform Engineering

Core Responsibilities: enable self-service, standardize runtimes/tooling, abstract infra, enforce compliance & security, provide observability.
8 Building Blocks (top→bottom): Self-Service Portal (IDP), Cost & Capacity, Observability, Policy & Security, Package & Artifact Mgmt, CI/CD Pipelines & Deployment, Infrastructure as Code, Compute & Orchestration.
Top Container Challenges (CNCF 2025): Cultural changes (47%), Lack of training (36%), Security (36%), CI/CD (35%), Monitoring (35%), Complexity (34%), Scaling, Networking, Logging, Service Mesh.

2 GitLab Pipelines

```
build-image:
  stage: release
  image: docker:28
  services:
    - name: docker:28-dind # Docker-in-Docker
      alias: docker
  variables:
    IMG: $CI_REGISTRY_IMAGE:latest
  before_script:
    - docker login -u $CI_REGISTRY_USER -p $CI_REGISTRY_PASSWORD $CI_REGISTRY
  script:
    - docker build -t $IMG -f Dockerfile.prod .
    - docker push $IMG
  rules:
    - changes: [Dockerfile.prod]
```

2.1 Core Concepts & DSL Mapping

Pipeline = automated steps → faster, safer, reliable. Traceable logs, parallel jobs, consistent envs kill “works on my machine”.
5-stage flow: commit/MR/tag → build (compile, image, package) → test (unit/integration/security/perf) → release (publish to registry) → deploy (target env) → post-deploy (smoke, health, rollback).

2.2 Jobs, Stages & Rules

.gitlab-ci.yml YAML DSL. **Job** = smallest unit, belongs to one **stage**, runs script on a runner. Same stage → parallel; stages → sequential.
rules: when a job runs (modern only/except). **needs:** DAG, start before stage finishes. **dependencies:** which jobs’ artifacts to pull (no ordering). **workflow:rules:** gate the whole pipeline.

```
workflow:
  rules:
    - if: $CI_PIPELINE_SOURCE == "merge_request_event"
    - if: $CI_COMMIT_BRANCH && $CI_OPEN_MERGE_REQUESTS
```

```
when: never
- if: $CI_COMMIT_BRANCH
deploy:
  stage: deploy
  needs: [build] # DAG, skip stage wait
  dependencies: [build] # pull build's artifacts
  rules:
    - if: '$CI_COMMIT_BRANCH == "main"'
    - when: never
```

2.3 Cache, Artifacts & Services

Cache: shared across jobs/pipelines, speed up repeat work (deps, compiled objects). **Artifacts:** files persisted to GitLab for downstream stages or download; declare via dependencies/needs. **Services:** sidecar containers (DB, dind, API) reachable via alias.

```
integration-test:
  services: [{name: postgres:17, alias: db}]
  variables: {POSTGRES_DB: app, POSTGRES_USER: u, POSTGRES_PASSWORD: p}
  script: [psql -h db -U $POSTGRES_USER -d $POSTGRES_DB -c "SELECT 1"]
  cache: {key: deps, paths: [".cache/"]}
  artifacts: {paths: [bin/], expire_in: 1h}
```

2.4 Environments

Where code is deployed. **Tiers:** Local (workstation) → Dev (sandbox) → Integration (CI target) → Test/QC (interface tests) → Staging (prod mirror) → Production (live). Linkable to K8s cluster (set up via UI). environment.on_stop unstages; action: stop stops it.

2.5 Push vs. Pull Deployment

Push (Jenkins/GitLab calls cluster): + easy, flexible targets (VMs/containers, not just K8s); – firewall open, admin access out, pipeline per env. **Pull** (agent in cluster, e.g. Argo/Flux): + secure (no inbound), auto-reconcile, scales identically; – agent per cluster, polling latency.

2.6 GitLab Runner

Worker executing jobs. **Scope:** Instance (shared, GitLab-managed) / Group / Project (dedicated). **Executors:** shell, ssh, docker, docker+machine, kubernetes, virtualbox, parallels. Register: gitlab-runner register with URL + token + tags + executor. Jobs pick runner by tags:.

2.7 More Keywords & Docker Build

default: fallback
image/before_script.
before_script/after_script wrap every script. rules:changes: run only if files changed.
Predefined vars:
\$CI_REGISTRY[_IMAGE],
\$CI_REGISTRY_USER/_PASSWORD,
\$CI_COMMIT_BRANCH,
\$CI_COMMIT_SHORT_SHA,
\$CI_PROJECT_DIR, \$CI_JOB_TOKEN (auth API/registry),
\$CI_MERGE_REQUEST_IID. Pin images

by digest (img@sha256:..), not :latest → flaky.

```
build-image:
  stage: release
  image: docker:28
  services:
    - {name: docker:28-dind, alias: docker} # Docker-in-Docker
  variables: {IMG: $CI_REGISTRY_IMAGE: $CI_COMMIT_SHORT_SHA}
  before_script:
    - docker login -u $CI_REGISTRY_USER -p $CI_REGISTRY_PASSWORD $CI_REGISTRY
  script:
    - docker build -t $IMG -f Dockerfile.prod .
    - docker push $IMG
  rules:
    - changes: [Dockerfile.prod]
```

3 Terraform

IaC: managing/provisioning infra through code instead of manual processes. **Declarative** (TF, what/desired state, implicit flow, stateful, idempotent) vs **Imperative** (shell/CLI, how/steps, explicit, stateless, not idempotent). TF talks to Provider plugins (RPC) → Client SDK (HTTP) → cloud API. **Resources** (resource) = managed objects. **Data sources** (data) = read-only lookups. HCL sample:

```
terraform {
  required_version = ">= 1.0.0"
  required_providers { aws = { source = "hashicorp/aws", version = "> 5.0" } }
}
provider "aws" { region = "us-east-1" }
variable "instance_type" { default = "t2.micro" }

resource "aws_instance" "web" {
  ami = data.aws_ami.ubuntu.id
  instance_type = var.instance_type
}
output "public_ip" { value = aws_instance.web.public_ip }
```

Workflow: write .tf → init (backend + download providers/modules) → plan → apply → destroy. Other: fmt, validate, console (REPL), state list/show <addr>. Reference attrs as <type>.<name>.<attr>.

3.1 OpenTofu

Fork of Terraform v1.5.7 under **Linux Foundation**, MPL-2.0 (TF switched to BUSL). Drop-in replacement, adds **state encryption**. Avoids single-vendor lock-in.

3.2 HCL Features

count (numeric loop, count.index) and **for_each** (map/set loop, each.key/value) for multiple instances. **Data source** example:

```
data "aws_ami" "ubuntu" {
  most_recent = true
  filter { name = "Name" values = ["ubuntu/images/hvm-ssd/ubuntu-*"] }
  filter { name = "virtualization-type" values = ["hvm"] }
  owners = ["099720109477"] # Canonical
}
```

3.3 Variables

Precedence: CLI -var/-var-file > env TF_VAR_<name> > *.tfvars > default. No default & not supplied → CLI

prompts operator. TF_VAR_ prefix ideal for secrets (TF_VAR_access_key=...). depends_on forces explicit ordering; implicit deps come from attribute references.

3.4 State

terraform.tfstate (JSON) maps HCL to real resource IDs (e.g. ARN). **3-way diff** on apply: **NEW** (HCL) vs **PREVIOUS** (.tfstate) vs **EXISTING** (cloud refresh) → decides actions, detects drift. **Remote backend** (S3) shares state in teams; **state locking** via use_lockfile (S3-native, replaces DynamoDB):

```
terraform {
  backend "s3" {
    bucket = "my-tfstate"
    key = "prod/terraform.tfstate"
    region = "us-east-1"
    use_lockfile = true
  }
}
```

3.5 File & Folder Layout

Files: terraform.tf (required_version/providers), backend.tf (remote state, kept out of VCS), providers.tf (auth/config), main.tf (resources/data), variables.tf (inputs, alphabetical), outputs.tf (returns, alphabetical), locals.tf (derived values), override.tf (loaded last, sparingly). **One folder per env** (dev/, prod/) → separate state files so a corruption hits only one env; shared modules/ along-side.

3.6 Modules

Reusable folder of .tf files (root config is itself a module):

```
module "vpc" { source = "../modules/vpc" cidr = "10.0.0.0/16" }
# reference output: module.vpc.vpc_id
```

3.7 Patterns & Recovery

import adopts existing infra; terraform apply -replace=<addr> (or legacy taint) forces a destroy+recreate of a single resource. Region-aware AMI via map:

```
import { to = aws_instance.web id = "i-0abc123" }
# then: apply

variable "AMIS" { type = map(string) default = { us-east-1 = "ami-05b1..." } }
ami = var.AMIS[var.region]
```

4 Ansible

Configures provisioned machines (install/configure software, deploy code); after provisioning. Idempotent (no-op if already in desired state), agentless, push-based, no statefile/locking. Needs Python on host & targets. Other CM tools: **Puppet**, **Chef**, **CFEngine**, **Juju**. **TF vs Ansible:** TF = provisioning, declarative, state, idempotent. Ansible = CM, hybrid (imperative tasks, declarative

time modules), no lifecycle, idempotency *optional*. Both: templating, agentless, OSS.

Execution: over SSH, copies generated Python module to ~/.ansible/tmp on target, runs with remote Python, returns JSON on stdout, deletes. Nothing stays installed. **Exceptions:** raw module runs cmd via SSH (no Python needed, e.g. bootstrap); connection: local runs on controller (network devices); Windows via **WinRM**.

4.1 Initial Setup & Inventory

ssh-keygen -t ed25519 -C *@y" on controller, ssh-copy-id user@server, test with ansible all -m ping. **Inventory** (INI or YAML) lists managed hosts grouped, referenced by play's hosts: . ansible-inventory -i inv --list verifies it.

```
[web]
web1.example.com
web2.example.com  ansible_host=10.0.0.5
[web:vars]
ansible_user=deploy
```

4.2 Playbooks, Plays & Tasks

YAML file = playbook (1..n plays); play = hosts + tasks; task = one module call. Run: ansible-playbook -i inv playbook.yml. Flags: --check (dry run), --list-tasks (preview), --tags/--skip-tags, --ask-vault-pass.

```
- name: Example
  hosts: web
  become: true # sudo
  gather_facts: false # skip setup (faster)
  force_handlers: true # run handlers even if a
    task fails
  vars:
    packages: [nginx, curl]
    secret: "{{ vault_password }}"
  roles:
    - myrole
  tasks:
    - name: Install packages # list-as-option >
      loop (faster, deps OK)
      ansible.builtin.apt: # FQCN: namespace.
        collection.module
        name: "{{ packages }}"
        state: present
        notify: restart nginx
    - name: Render config
      ansible.builtin.template: { src: app.conf,
        j2, dest: /etc/app.conf }
    when: enable_service
    tags: config
  handlers:
    - name: restart nginx
      ansible.builtin.service: { name: nginx,
        state: restarted }
```

4.3 Ad-hoc, Modules & Errors

ansible -i inv -m <mod> -a '<args>' <target>. setup = gather facts (ansible facts). command/shell/raw not idempotent; shell allows pipes/vars. Playbook stops at first fail unless ignore_errors: true (use sparingly). **Handlers:** only run when notified and all prior tasks OK; force_handlers overrides.

```
- name: validate
  ansible.builtin.command: httpd -t
  register: out
```

```
changed_when: out.failed
failed_when: false
- ansible.builtin.file: { path: /data, state:
  directory } # file|link|touch|absent
- ansible.builtin.lineinfile: { path: /etc/x.
  conf, line: "Include sites/*" }
```

4.4 Vault, Collections, Roles & Tags

Vault: ansible-vault create|edit|view|encrypt foo.yml; encrypt_string 'secret' --name 'pw' inlines a single var. Pass via --vault-id @prompt|/path/file|name@script. Ciphertext safe to commit. **Collections** bundle modules/roles/plugins. FQCN ns.coll.mod (e.g. cisco.ios.ios_banner); independent versioning. Install: ansible-galaxy collection install <n> or via requirements.yml (collections: + roles:) then install -r. Sources: **Galaxy** (community), **Automation Hub** (RedHat, certified). **Roles:** ansible-galaxy init <n> creates dirs tasks/ handlers/ defaults/ vars/ files/ templates/ meta/ tests/; use via roles: list. **Tags:** --tags/--skip-tags. Inherit from play/import_* (static); include_* (dynamic) needs explicit apply: tags: . Specials: *always* runs unless --skip-tags=always; *never* runs only via --tags=never.

4.5 Jinja2 Templating

Rendered on **controller**, only result sent to target. {{ inventory_hostname }} = current target. Filters: ipaddr/ipv4/ipv6 (validate/version), hash('sha1'), default(x), min/max/shuffle.

```
{% for host in groups['webservers'] %}
BalancerMember http://{{ host }}:{{ port |
  default(80) }}
{% endfor %}
```

5 Kubernetes (K8)

Object: persistent record-of-intent. **Controller:** reconciles current → desired state.

5.1 Architecture

Control plane: kube-apiserver (REST front), etcd (KV store), kube-scheduler (pod → node), kube-controller-manager (built-in controllers), cloud-controller-manager (cloud glue, optional). **Node:** kubelet (runs pods via container runtime), kube-proxy (Service networking/iptables).

5.2 Objects

Pod: smallest unit, 1+ containers share net/storage. **Sidecar:** aux container alongside main (logs, proxy).

Init container: runs to completion before main. **ReplicaSet:** keeps N pod replicas (rarely used directly). **Deployment:** manages ReplicaSets, rolling updates, rollback. **StatefulSet:** ordered pods with stable identity + per-pod PVC (DBs, brokers). **DaemonSet:** one pod per (matching) node - logs, CNI, monitoring. **Service:** stable VIP/DNS + LB over pod set. **Ingress:** L7 HTTP(S) endpoint. **Volume:** dir mounted into pod. **ConfigMap:** non-secret KV config. **Secret:** sensitive KV, base64-encoded, *not* encrypted by default. **HPA:** autoscales replicas on metrics.

```
apiVersion: apps/v1
kind: Deployment
metadata: { name: app }
spec:
  replicas: 3
  selector: { matchLabels: { app: web } }
  strategy:
    type: RollingUpdate
    rollingUpdate: { maxUnavailable: 1, maxSurge:
      1 }
  template:
    metadata: { labels: { app: web } }
    spec:
      initContainers:
        - { name: init, image: busybox, command: [
          'sh', '-c', 'sleep 10'] }
      containers:
        - name: web
          image: nginx
          ports: [{ containerPort: 80 }]
        - { name: sidecar, image: busybox, command:
          ['sh', '-c', 'while true; do sleep 30;
            done'] }
```

5.3 Services

Stable VIP + DNS over label-selected pods, round-robin LB. **ClusterIP** (default) internal only; **NodePort** static port on every node <NodeIP>:<port>; **LoadBalancer** cloud LB w/ public IP; **ExternalName** CNAME alias to external DNS (no proxy). **Ingress** routes HTTP(S) by host/path (needs controller: Nginx, Traefik). **Gateway API** (Gateway+HTTPRoute/TLSRoute) - modern, multi-protocol, role-oriented Ingress replacement. *Port trap:* Service port vs targetPort (container) vs nodePort; Ingress backend refs Service port.

5.4 ConfigMaps & Secrets

Inject single var (configMapKeyRef/secretKeyRef), all keys (envFrom), or mount as files.

```
env:
- name: APIKEY
  valueFrom: { configMapKeyRef: { name: cfg, key:
    api-key } }
envFrom:
- configMapRef: { name: cfg }
volumeMounts:
- { name: v, mountPath: /etc/cfg }
volumes:
- { name: v, configMap: { name: cfg } } # or:
  secret: { secretName: s }
```

5.5 Storage: PV, PVC, StorageClass

Container FS is ephemeral. **Volume:** dir in pod, lives w/ pod (emptyDir, hostPath, NFS, ...). **PersistentVolume (PV):** cluster-scoped piece of

external storage. **PersistentVolumeClaim (PVC):** namespaced request for storage (size, accessMode); binds to a matching PV. **StorageClass:** dynamic provisioner - creates a PV on-demand if no static PV matches (e.g. gp3, standard-rwo).

5.6 LimitRange & ResourceQuota

LimitRange: per-container min/max + default req/lim in NS; out-of-range Pod rejected, missing req/lim auto-injected. **ResourceQuota:** caps summed NS usage; over-quota Pod rejected.

5.7 Namespaces

Logical isolation in one cluster (RBAC, quotas). Namespaced (Pods, Services) isolated; cluster-wide (Nodes, PV) shared. Defaults: default, kube-system, kube-public, kube-node-lease.

5.8 Rolling Updates

Deployment ensures availability during change. **Recreate:** kill all → recreate (downtime). **RollingUpdate** (default): tuned by maxUnavailable (max pods down) + maxSurge (max extra pods). kubectl rollout history/undo for revisions.

5.9 Scheduling

kube-scheduler picks node via **Filtering** then **Scoring**. *Filtering* (feasibility): PodFitsHostPorts, PodFitsResources, PodMatchNodeSelector, CheckNodeDiskPressure, CheckVolumeBinding. *Scoring* (rank): SelectorSpreadPriority (spread same Service/STS/RS), LeastRequestedPriority (free resources), NodeAffinityPriority.

```
spec:
  nodeSelector: { disktype: ssd }
```

Taints on nodes repel pods; **Tolerations** on pods allow (don't force) them onto tainted nodes - pair w/ nodeSelector/affinity to pin. Effects: *NoSchedule*, *PreferNoSchedule*, *NoExecute* (evict). No effect on DaemonSets.

CLI: kubectl taint nodes <n> gpu=true:NoSchedule, kubectl label nodes <n> disktype=ssd (trailing - removes).

```
tolerations:
- { key: "gpu", operator: "Equal", value: "true",
  effect: "NoSchedule" }
```

cordon/drain: kubectl cordon <n> marks unschedulable (existing stay); drain <n> also evicts pods (maintenance); uncordon <n> reverts.

5.10 Probes & Resources

kubelet probes via httpGet/tcpSocket/exec. **Liveness:** fail → restart. **Readiness:** fail →

remove from Service endpoints. **Startup:** guards slow boots, disables others until OK. Repeated liveness fail → exp backoff → CrashLoopBackOff. **requests:** reserved min (scheduling). **limits:** ceiling - CPU throttled, memory over-limit → OOMKilled.

```
resources:
  requests: { cpu: "250m", memory: "64Mi" }
  limits: { cpu: "500m", memory: "128Mi" }
livenessProbe:
  httpGet: { path: /healthz, port: 80 }
  initialDelaySeconds: 5
  periodSeconds: 10
readinessProbe:
  httpGet: { path: /ready, port: 80 }
```

5.11 Commands

Apply: kubectl {create|apply|replace} -f manifest.yaml
Gen YAML: kubectl create deploy web --image=nginx --dry-run=client -o yaml > d.yaml
Imperative: kubectl run, scale deploy x --replicas=4, expose Update image: kubectl set image deploy/x web=nginx:1.27
Autoscale: kubectl autoscale deploy x --cpu-percent=50 --min=1 --max=10
Inspect: kubectl get pods -A -o wide, describe node <n>, get nodes --show-labels
Exec: kubectl exec -it nginx-xxx -- sh (drop -it for one-shot)
Rollout: kubectl rollout {status|history} deploy x, rollout undo deploy x --to-revision=N
Node ops: kubectl cordon|uncordon|drain <n>
Default ns: kubectl config set-context --current --namespace=lab

6 Helm

Package manager for K8 (apt/brew-like). **Chart**=files describing related K8 resources. **Release**=running instance. **Repo**=stores versioned charts. **Arch:** client-side only, talks K8 API direct, release state in K8 **Secrets** (no Tiller). **Pros:** huge chart library, versioning/deps, lifecycle (install/upgrade/rollback), powerful Go-template logic. **Cons:** templates hard to debug, ßpaghetti YAML, opaque state, steep curve. Use for: distributing complex apps, consuming community charts, advanced logic. Avoid for: simple YAML tweaks.

6.1 Commands

```
helm repo add|list|rm <repo> <url> # repos
helm search repo <r> <c> -1 # all versions
helm install -f vals.yaml <rel> <chart> --
  version 1.2.3
helm upgrade <rel> <chart>; helm rollback <rel>
  <rev>
helm uninstall <rel> # also drops history
helm list; helm history <rel> # CHART vs APP
VERSION
helm status <rel> # unknown|deployed|uninstalled|
```



```
# superseded|failed|uninstalling|pending-|
install,upgrade,rollback}
helm get all|values|notes|manifest <rel> #
values=overrides only
helm template <n> <c> -f v.yaml --set k=v #
render local
helm install <n> <c> --dry-run --debug #
simulate
helm lint <chart> [--with-subcharts]
helm show chart|readme|values|all <chart>
helm package <dir> # build .tgz; helm push <pkg>
<repo>
helm dependency update # pulls into charts/
```

6.2 Chart Structure & Chart.yaml

helm create <n> scaffolds:
Chart.yaml (meta), values.yaml (defaults), templates/ (Go-tmpl manifests), templates/_helpers.tpl (named templates, _=not rendered), templates/NOTES.txt (rendered post-install), charts/ (subcharts), crds/, .helmignore (globs, !=negate; excludes from helm package).
Chart.yaml fields: **apiVersion** (v2=Helm 3+, v1=legacy), **name**, **version** (SemVer2; pkg name = name+ver), **type** (application default | library = helpers only, no install), **appVersion** (app being shipped, info-only), **kubeVersion** (range like >=1.12.1 <1.14.0 || >=1.14.2; install fails if mismatch), **dependencies** (name/version/repository/condition/alias; requirements.yaml deprecated, lockfile on update).

6.3 Templating

Top-level objects:
.Release. {Name, Namespace, Revision, .Chart.*, .Values, .Files, Get <path>, .Capabilities. {APIVersions, KubeVersions, Template. {Name, BasePath}.
Values precedence (low→high, last wins): subchart values.yaml → parent values.yaml → -f myvals.yaml → --set k=v. Unset default: --set x=func.
Functions: Sprig lib (~60+: upper quote eq default trunc toYaml indent nindent). Pipelines | pass result as last arg. include = template but pipeable. required mmsg".Values.x fails if empty. tpl <str> . evaluates string as template.
Flow: {{ if/else if/else }}...{{ end }} (false on 0//nil/empty); {{ range \$k,\$v := .Values.list }}; {{ with .Values.sub }} rescopes .. Vars: {{- \$n := .Release.Name -}}.
Whitespaces: {{- strips left, -}} strips right; nindent N = newline + indent. YAML indent critical.

```
{{- define "app.name" -}}{{ default .Chart.Name .
Values.name | trunc 63 }}{{- end }}
name: {{ include "app.name" . | quote }}
{{- required "set .Values.img!" .Values.img }}
resources: {{- toYaml .Values.resources |
nindent 4 }}
```

7 Kustomize

Template-free overlay tool, built into kubectl (-k). kustomization.yaml lists **resources** + declarative **transformers** (cross-cutting fields) + **patches** (partial spec edits) + **generators**. **Pros:** pure YAML, native to kubectl, easy diff/review, easy to learn. **Cons:** less powerful than Helm, no deps mgmt, hard to package/distribute, limited conditionals. Use for: env-specific configs (dev/prod), avoiding templating. Often combined with Helm.

7.1 Commands

```
kustomize create --resources deploy.yaml,svc.
yaml # scaffold
kustomize edit add resource|set image|add label|
remove resource # render to stdout
kustomize build <dir> | kubectl apply -f -
kubectl apply -k <dir> # native
kubectl kustomize <dir> # render only (no apply)
kubectl delete -k <dir>
kustomize cfg fmt <dir> # format yaml
kustomize localize <dir> # bundle remote URLs|git
locally
kustomize edit fix # auto-migrate deprecated
syntax
```

7.2 kustomization.yaml

Deprecated → **new:**
commonLabels → labels,
patchesStrategicMerge/patchesJson6902 → patches, bases: → inside resources:.

```
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization
resources: [../base, storage.yaml] # files,
dirs, URLs
namespace: prod # transformers
labels: {app:nginx,upgrade:Service},
nameSuffix: -v1
commonAnnotations: {version: 1.0.0}
labels: # replaces commonLabels
+ includeSelectors: true
bases: {app:kubernetes.io/name: hello}
images: # change tag/name w/o editing Deployment
- {name: nginx, newName: bitnami/nginx, newTag: latest}
patches: # strategic or JSON6902
- target: {kind: Service, name: nginx-app}
patch: |-
- {op: replace, path: /metadata/name, value:
nginx-server}
configMapGenerator: # hash suffix appended ->
change forces rolling restart
- name: app-config
literals: [LOG_LEVEL=debug]
files: [config.env] # or files:
secretGenerator: # base64-encodes; same hash
mechanism
- {name: ex-secret, literals: [user=admin, pass=
secret]}
components: [../components/mysql]
```

ConfigMap/Secret modes: (1) *Traditional:* declare as resource → change does **not** restart pods (silent drift). (2) *Generator:* configMapGenerator/secretGenerator appends content hash to name → references rewritten → pods rolling-restart on change.

7.3 Bases, Overlays, Components

Base: common (base/ with kustomization.yaml). **Overlay:** env-specific dir (overlays/dev, overlays/prod) that lists base in

resources: + adds patches/transformers → produces a **variant**.
Component: kind: Component (v1alpha1); reusable mix-in chunk (resources+patches+generators), included via components: in overlays. resources = full manifests/bases; bases: key itself is deprecated.

```
# components/mysql/kustomization.yaml
apiVersion: kustomize.config.k8s.io/v1alpha1
kind: Component
resources: [deployment.yaml, service.yaml]
```

7.4 Best Practices & Pitfalls

Keep common values (namespace, labels) in base. Organize by kind (services.yaml, hpa.yaml). Standard dirs: base/, overlays/, components/, patches/.
Pitfalls: nested overlays (flatten!); not pinning image tags (use images: with explicit newTag); drift between cluster + kustomization (no release state tracked, unlike Helm); skipping kustomize build validation before apply; using traditional ConfigMaps when rollout-on-change needed (use generator).

7.5 Helm vs. Kustomize

Need
Distribute complex app (DB+App)
Env-specific configs (dev/prod)
Avoid new templating language
Consume community charts
Loops/conditionals/deps

8 K8s Monitoring & Logging

Observability = property (degree system generates actionable insights, qualitative): **Observing/Monitoring** = process (collect/store/evaluate *signals*). **Three pillars:** *Metrics* (quantitative: util, counters, rates), *Logs* (qualitative: stack traces, events), *Traces* (request path across services via *spans* – latency/bottlenecks, e.g. Tempo, Jaeger). **Monitoring vs Logging** complement each other.

K8s obs gaps: no historical metrics, kubectl logs only running containers (lost when container dies, kubelet cleanup), no native dashboards/alerts. Container logs at /var/log/pods/...; kubelet rotates at 10MB, keeps 5.
Push (agent-based): agent on system sends to collector; observed decides when. **Pull (agentless):** collector queries via SNMP/ICMP/HTTP; collector decides, but system must be online. **Golden Signals** (Google SRE): latency, traffic, errors, saturation.

PLG stack (lightweight, K8s-native): Prometheus+Loki+Grafana. **LGTM:** Loki, Grafana, Tempo, Mimir. **EFK/ELK:** Elasticsearch (Java, JSON index, full-text, heavy) + Kibana + Fluentd/Logstash. PLG indexes *labels* only (cheap, fast, weaker text search). **OpenTelemetry:** unified telemetry (metrics/logs/traces) spec.

8.1 Prometheus & PromQL

Go, TSDB, pull-based scrape over HTTP (/metrics). Pushgateway/remote-write for short jobs. Client libs Go/Java/Python/Ruby/Rust. Standalone (no remote dep).
Targets: kubelets, kube-state-metrics, exporters. **Exporters:** bridge for apps without native /metrics (PostgreSQL, NGINX, Redis, JIRA, Kafka, ...); use when can't instrument. **node-exporter:** HW/OS. **kube-state-metrics:** K8s object state. Install metrics-server for kubectl top. CPU units: 100m=0.1 CPU absolute.
Metric types: *Counter* (monotonic, reset to 0: requests/errors), *Gauge* (current value: memory, concurrency), *Histogram* (sampled buckets), *Summary* (sampled + total/sum). **Format:** name{label=val,...} value
Helm → foundation of **OpenMetrics**
PromQL data types: *Instant vector* (set of series, one sample each, same ts), *Range vector* (range per series, e.g. [5m]), *Scalar* (float), *String* (unused). **Offset** (positive = backwards): rate(x[5m] offset 1w). **@-modifier:** sum(x 1609746000). histogram_quantile(0.95, rate(..._bucket[5m])) for p95 from Histogram. rate() useless on spikes/outliers – use raw & > threshold.

```
sum by (namespace) (kube_pod_status_ready{
condition="false"})
rate(mysql_global_status_commands_total{command=
"select"}[5m])*100 > 35
```

Prometheus Operator (kube-prometheus-stack): bundles operator+Prom+Grafana+Alertmanager. CRDs: Prometheus/ServiceMonitor/PodMonitor. Match chain: Prom serviceMonitorSelector→SM selector→Service selector→Pods. **Scaling Prometheus:** *Federation* (hierarchical: region scrapes DC); *Thanos* (sidecar, global query, long-term, dedup); *Grafana Mimir* (horiz. scalable, ingester/querier components, multi-tenant, remote-write only).

8.2 Grafana, Alerting

General-purpose viz, many data sources (Prom/Loki/ES). *Explore* feature: browse data sources, build queries, debug, share. OAuth/SAML/LDAP. Dashboard JSON repo grafana.com/grafana/dashboards. Tips: important data top, status-light > number, simplicity, Data Source variable.

Alertmanager vs Grafana Alerts: Alertmanager (Prom ecosystem, GitOps-friendly, separates viz/alert, Prom-only). Grafana Alerts (any data source, tied to dashboards, can chain to ext Alertmanager, less GitOps-friendly). Flow: *Rule*→evaluate/threshold→*Alert*→Silence?–(Email/Slack/Teams/Discord/Threema/WhatsApp Points). **Warning ≠ Alert:** cert expires in 45d, CPU 75%, budget 200% are warnings not pages. Avoid alert fatigue, build up slow, test.

8.3 Logging: Loki, Alloy

Architectures: *Node-level* (DaemonSet agent reads node log file – survives container death, batched); *App-level* (app pushes to backend directly); *Sidecar* (streaming = sidecar writes to file then node agent ships – kubectl logs won't show sidecar's own; logging-agent sidecar = ships to backend directly, low-latency but loses on disconnect, adds complexity, misses crashes). Best practice: structured JSON logs.

Loki: horiz. scalable, indexes *labels* only (not text), *nanosecond* unix ts, cheap storage, PromQL-like **LogQL:** {job=ttest"} != "error"; metric queries rate({app=*}) != "176"[1m]); filters != contains, !=, != regex, !=.

Alloy (Grafana, replaces Promtail): collects all telemetry, own HCL-like lang. Components: loki.source.file (tail files, classic DaemonSet pattern), loki.source.kubernetes (subscribe via K8s API – no root, no DaemonSet needed, but more CPU/net), loki.source.kubernetes_events (K8s events). Logs enriched with **K8s metadata**, forwarded over HTTP. Helm-installable.

Pitfalls: too much logging → storage/complexity; no alert tuning → noise/missed; forgetting rotation/archival; relying on single signal; logging secrets.

9 Service Mesh

9.1 Microservices

Small independent services, one business function, communicate via lightweight APIs. Microservices == distributed computing.

Pros: independent deploy, flexibility, scalability, fault isolation, team auto-

nomy.
Cons: operational complexity, distributed-system challenges, data consistency, higher resource use.
Challenges: *Communication* - services must know endpoints, new service ripples changes. *Security* - in-cluster traffic unsecured, any pod → any pod. *Monitoring* - spread out, non-business logic added per app.

9.2 Distributed Computing Fallacies

Network reliable; latency zero; bandwidth infinite; topology stable; one admin; transport cost zero; network homogeneous. All false → motivate mesh.

9.3 Service Mesh

Dedicated infra layer handling service-to-service communication: traffic routing, security, observability, resiliency - abstracted from app code.

Capabilities: Traffic Mgmt (routing, splitting), Resilience (fault tolerance, retries, fallback), Security (mTLS, authN/Z), Observability (metrics, logs, tracing).

Why: centralized policy, uniform observability, faster incident resolution, transparent security, advanced rollouts (canary/AB), resilience.

Components: *Data Plane* - lightweight proxies (Envoy), handle comm/metrics/security. *Control Plane* - manages config/policies (Istiod), pushes rules to data plane.

9.4 Mesh Evolution: Library → Sidecar → Kernel

Library: mesh logic linked into each app, per-language lib (Python, Go, ...). Invasive, language-locked.

Sidecar: proxy per pod, transparent traffic interception, no code change, language-agnostic. Downsides: invasive pod-spec mutation, upgrade = pod restart, large CPU/RAM per pod (worst-case provisioning), extra hop.


Kernel/Ambient: mesh in kernel (eBPF) or per-node agent. 30 pods/node → 1 proxy/node instead of 30. L7 still typically needs userspace proxy.

Leading meshes: *Istio* (Google, feature-rich), *Linkerd* (lightweight, perf), *Kuma* (CNCF, VMs+containers), *Consul Connect* (HashiCorp, discovery), *Cilium* (eBPF, not dedicated mesh).

9.5 Sidecar Pattern Pros/Cons

Pros: + transparent (no code change, lang-agnostic), + per-pod isolation, + platform team owns cross-cutting (mTLS/retries/routing/telemetry).

Cons: - CPU/RAM scales with pod count, - extra hop = latency, - sidecar injection + proxy upgrade needs restart, - overkill for few services.



Envoy proxies authenticate via **mTLS** (mutual TLS, both sides present certs).

9.6 Ambient Mesh (Sidecarless Istio)

Sidecar downsides → ambient: per-node **zTunnel** agent (L4, mTLS, authN/Z, no HTTP parsing) communicates across nodes via HTTPS CONNECT tunnel. Optional L7 **Waypoint Proxy** per namespace (Envoy) for rich policies/telemetry. zTunnel alone = L4 metrics only.

Slicing: Secure Overlay (L4): TCP routing, mTLS tunneling, simple authZ, TCP metrics. **L7 Processing:** HTTP routing/LB, circuit breaking, rate limiting, fault injection, retry, timeouts, rich authZ, HTTP metrics/access logs/tracing.

Cilium variant: dynamic **eBPF** programs in kernel handle L3/L4 (network path); per-node Envoy proxy for L7 termination when needed. No sidecar, no extra hop for L4.

9.7 Gateway API & Ambient Routing / GAMMA

Enable ambient: `kubectl label ns x istio.io/dataplane-mode=ambient (sidecar mode: istio-injection=enabled)`. Routes via K8s **Gateway API** (HTTPRoute) not VirtualService. Two Deployments behind one Service ≈50/50 default; override via weights. Best practice: version label per Deployment.

GAMMA (Gateway API for Mesh Management & Administration): Gateway API for east/west (mesh) traffic, not just north/south (ingress). Adopted by more meshes.

```
# weighted canary 90/10:
- backendRefs:
  - { name: frontend, port: 5000, weight: 90 }
  - { name: frontend-v2, port: 5000, weight: 10 }
# mirror/shadow (copy to v2, drop response):
- backendRefs:
  - { name: frontend, port: 5000 }
  filters:
  - type: RequestMirror
    requestMirror: { backendRef: { name: frontend
      ~v2, port: 5000 } }
```

istio-system pods: Istiod, Ingress Gateway (Envoy), Prometheus, Grafana, Jaeger, Kiali. **Kiali shapes:** rectangle=app, triangle=Service, square=workload/pod.

9.8 Traffic Management (Istio Classic)

Gateway replaces K8s ingress/egress. VirtualService binds gateway, routes

to services. DestinationRule applies post-routing policy: LB, connection pool, outlier detection (circuit breaking), TLS mode, defines subsets (versions).

```
apiVersion: networking.istio.io/v1beta1
kind: Gateway
metadata: { name: my-gateway }
spec:
  selector: { istio: ingressgateway }
  servers:
  - port: { number: 80, name: http, protocol: HTTP }
    hosts: ["example.com"]
  --
kind: VirtualService
metadata: { name: my-vs }
spec:
  hosts: ["example.ch"] # must match gateway host
  gateways: [my-gateway]
  http:
  - route:
    - destination:
        host: my-service
        weight: 99 # split traffic if 2nd dest defined
      port: { number: 80 }
  --
kind: DestinationRule
metadata: { name: my-dr }
spec:
  host: my-service
  trafficPolicy:
    tls: { mode: ISTIO_MUTUAL }
```

9.9 Resilience & Use Cases

Fault Injection: simulate net failures, test error handling. **Timeouts:** cap request time, avoid queue buildup. **Retries:** auto-retry on failure.

Cascading Failure: one service dies → callers pile up → whole graph red.

Circuit Breaking: trip on errors/outlier → prevent cascading failure (DestinationRule outlierDetection).

Rate Limiting: cap RPS, protect downstream from overload.

Zero-Trust: encrypt + authenticate every call, no implicit trust in cluster. Istio exports Prometheus metrics by default.

9.10 Deployment Strategies

Rolling: gradual pod replace, default K8s. **Recreate / Big-bang:** kill all, deploy new (downtime). **Canary:** small % to v2, watch, ramp up. **Blue/Green:** two envs, switch traffic atomically, easy rollback. **Dark Launching / Shadow:** mirror prod traffic to v2, drop response (test under real load). See *General Definitions*.

9.11 Security

AuthN: mTLS between services. **AuthZ:** AuthorizationPolicy = fine-grained RBAC.

```
# enforce mTLS in namespace
kind: PeerAuthentication
metadata: { name: default, namespace: production }
spec:
  mtls: { mode: STRICT }
  --
kind: AuthorizationPolicy
metadata: { name: require-jwt, namespace: foo }
spec:
  selector: { matchLabels: { app: httpbin } }
  action: ALLOW
  rules:
```

```
- from: [{ source: { namespaces: ["frontend"] } }]
to: [{ operation: { methods: ["GET"], paths: ["/headers"] } } ]
```

10 Cilium & eBPF

10.1 eBPF

eBPF (extended Berkeley Packet Filter): runs sandboxed programs directly in the Linux kernel without changing kernel source or loading modules. Programs attach to hooks (syscalls, network events) and are verified before loading. Enables high-performance networking, observability and security at kernel level.

10.2 Cilium

A CNI (Container Network Interface) plugin built on eBPF. Replaces kube-proxy: instead of one iptables rule chain per Service, it programs eBPF maps in the kernel (faster, scales better).

Identity-based security: Cilium derives a security *identity* from a pod's labels (not from its IP). Policy is enforced on identities, so it survives pod rescheduling and IP changes.

Hubble: Cilium's observability layer; gives network flow visibility (service map, L3-L7 flows).

10.3 Network Policies

Declarative firewall rules for pod traffic. Plain NetworkPolicy works at L3/L4 (IP, port, protocol). CiliumNetworkPolicy adds L7 (HTTP method/path, DNS, gRPC, Kafka).

Default deny: as soon as a pod is selected by any policy, all non-matching traffic is dropped (whitelist model). Rules split into **ingress** (incoming) and **egress** (outgoing) and are *stateful*: return traffic of an allowed connection is permitted automatically.

```
apiVersion: cilium.io/v2
kind: CiliumNetworkPolicy
metadata:
  name: allow-frontend
spec:
  endpointSelector:
    matchLabels:
      app: backend # applies to backend pods
  ingress:
  - fromEndpoints:
    - matchLabels:
        app: frontend # only frontend may connect
  toPorts:
  - ports:
    - port: "80"
      protocol: TCP
  rules:
  - http: # L7 rule
    - method: GET
      path: "/api/*"
```

11 GitOps

Git = single source of truth for infra/app config; in-cluster agent reconciles live → desired state.

4 Principles (OpenGitOps): *declarative* state; *versioned* & *immutable* (history = audit, revert = rollback);

pulled by agents; *continuously reconciled* (detect + self-heal drift).

11.1 Push vs. Pull

Push (classic CI/CD): pipeline holds creds, runs kubectl apply. + fast, well-known; - creds outside cluster, inbound firewall, no drift detection, perv-enp pipeline edits.

Pull (GitOps): in-cluster agent watches Git, syncs. + no inbound access, drift correction, multi-cluster scale; - agent per cluster, polling latency. *Trust:* CI server → agent + Git. *Access dir:* push = in; pull = out.

11.2 Argo CD vs. Flux CD

Argo CD: Web UI + CLI, multi-cluster from one UI, SSO/RBAC, sync waves, pre/post hooks, health checks. CRDs: Application (repo+path → cluster/ns + sync-Policy), ApplicationSet (templated per branch/cluster/env), AppProject (restrict what/where/who). Pattern: *App-of-Apps* (root App syncs children).

Flux CD: CLI-first, no UI, modular *GitOps Toolkit* (source, customize, helm, notification, image). Supports Gitless GitOps (OCI/S3) + image automation. *Naming clash:* `kustomize.config.k8s.io` (feature) vs. `kustomize.toolkit.fluxcd.io` (Flux CRD).

11.3 Secrets & Best Practices

Encrypted in Git: Sealed Secrets (controller decrypts), SOPS. **Referenced:** External Secrets, Secrets Store CSI driver (mounts Vault/KMS as volume).

Monorepo: clusters/<c>/bootstrap, apps/base/ shared, apps/<c>/ per-cluster. **Rules:** never commit plaintext, pin image tags (no latest), monitor controller, split monolith repos.

12 Site Reliability Engineering

"Keep the site up, whatever it takes"

SL Indicators: A measurement (e.g. Response latency over a period of 10 min) **SL Objectives:** A Objective (e.g. 99.9% below 5ms) **SL Agreements:** A economic incentive (e.g. or less bonus) **Error Budget:** 1 - SLO (e.g. 100% - 99.9% = 0.1%)

The four golden signals: latency, traffic, errors, and saturation.

12.1 Outage Handling

Goals: 1. Minimize impact (Automated, fast detection; good diagnostic information, practice remediation) 2. Prevent reoccurrence (handle event, write post-mortem)